



Partial-Update Decoding: recompute only $U_t \cup U_{t-1}$, reuse cached deltas

Transformer Block (expanded — AdaLN-Zero / DiT-style)

AdaLN-Zero: $\text{cond } c \rightarrow \text{SiLU} \rightarrow \text{Linear} \rightarrow (\gamma_1, \beta_1, \alpha_1 \mid \gamma_2, \beta_2, \alpha_2)$ [γ, β : modulate $\cdot \alpha$: gate each sublayer]



Block: $x \leftarrow x + \alpha \cdot \text{sublayer}((1 + \beta) \cdot \text{RMSNorm}(x) + \gamma)$ • SwiGLU FFN: $W_2(\text{SiLU}(W_1 x) \odot W_3 x)$ • Attn: $W_q, W_k, W_v \rightarrow \text{QK-RMSNorm} \rightarrow \text{SDPA} \rightarrow W_o$



U_t — newly unmasked at step t (recomputed)



U_{t-1} — unmasked at step $t-1$ (still in active set)



inactive — not recomputed, reuse cached delta